

CS 199
HW2 - Predicting Crime Rates
April 8, 2014

Sam Laane <laane2@illinois.edu>
José Vicente Ruiz <ruizcep2@illinois.edu>

Contents

1	Removing variables with missing values	2
1.1	Implementation	2
2	Basics - Linear regression	2
2.1	Implementation	2
2.2	Results	3
2.3	Conclusions	3
3	Basics - Nearest Neighbour regression	3
3.1	Implementation	3
3.2	Results	4
3.3	Conclusions	4
4	Dealing with missing values	4
4.1	Implementation	4
4.2	Results	5
4.3	Conclusions	5
5	Modified Nearest Neighbours	5
5.1	Implementation	5
5.2	Results	5
5.3	Conclusions	5

1 Removing variables with missing values

1.1 Implementation

```
1 # Read the data.
2 data <- read.csv('communitiesH.data', h=T) # Headers = True.
3
4 # Remove non-predictive attributes (including state number).
5 clean_data <- data[, -c(1:5)]
6
7 # Remove the attributes with question marks (unknowns) of them.
8 unkw <- clean_data == '?'
9 unkw_per_attr <- apply(unkw, 2, sum) # Sum the columns (2).
10 attr_with_unkw <- which(unkw_per_attr > 0) # Equivalent to find in
    Matlab.
11 no_unkw_data <- clean_data[, -attr_with_unkw]
```

2 Basics - Linear regression

2.1 Implementation

```
1 # Compute the regression with the data free of unknowns.
2 linear_regr <- lm(ViolentCrimesPerPop ~ ., data=no_unkw_data)
3
4 # Evaluate the regression looking at the mean-squared error on the
    training data.
5 residuals <- resid(linear_regr) # Function to get the residuals.
6 plot(predict(linear_regr, no_unkw_data), residuals)
7
8 mse_residuals <- sum((residuals - mean(residuals)) ^ 2) / length(
    residuals)
9 printf(" - Mean-squared error on the whole data: %.2e", mse_residuals)
10
11 # Split off some test data and compute a new regression on the rest of
    the data
12 # (train data).
13 folds <- cvFolds(nrow(no_unkw_data), K=5)
14 train_data <- no_unkw_data[folds$subsets[folds$which != 1], ]
15 test_data <- no_unkw_data[folds$subsets[folds$which == 1], ]
16
17 linear_regr_test <- lm(ViolentCrimesPerPop ~ ., data=train_data)
18
19 # Evaluate the regression looking at the mean-squared error on that
    test data.
20 residuals_test <- test_data$ViolentCrimesPerPop - predict(linear_regr_
    test, test_data)
21 plot(test_data$ViolentCrimesPerPop, residuals_test)
22
23 mse_residuals_test <- sum((residuals_test - mean(residuals_test)) ^ 2)
    / length(residuals_test)
24 printf(" - Mean-squared error on the test data (20%%): %.2e", mse_
    residuals_test)
25
```

```

26 # Prepare the data for the Box-Cox tranformation substituting zero
    values by
27 # a very small number.
28 boxcox_data <- no_unkw_data
29 boxcox_data$ViolentCrimesPerPop[ which(boxcox_data$ViolentCrimesPerPop
    == 0) ] <- 1e-100
30
31 # Apply Box-Cox and get a list of the lambdas and their log likelihood
    and obtain
32 # the value with the highest likelihood.
33 linear_regr_mod <- lm(ViolentCrimesPerPop ~ ., data=boxcox_data)
34 lambdas <- boxcox(linear_regr_mod, plotit=F)
35 max_lambda <- lambdas$x[ which(lambdas$y == max(lambdas$y)) ]
36
37 # Change the very low values to zero again.
38 boxcox_data$ViolentCrimesPerPop[which(boxcox_data$ViolentCrimesPerPop
    == 1e-100)] <- 0
39
40 # Transform the data with the given value of lambda.
41 if (max_lambda == 0) {
42   boxcox_data$ViolentCrimesPerPop <- log(boxcox_data$
    ViolentCrimesPerPop)
43 } else {
44   boxcox_data$ViolentCrimesPerPop <- (boxcox_data$ViolentCrimesPerPop^
    max_lambda - 1)/max_lambda
45 }
46
47 # Compute the linear regression again and plot it.
48 boxcox_regr <- lm(ViolentCrimesPerPop ~ ., data=boxcox_data)
49 plot(boxcox_regr)

```

2.2 Results

2.3 Conclusions

3 Basics - Nearest Neighbour regression

3.1 Implementation

```

1 # Compute the regression and plot it.
2 nn_regr <- knn.reg(train_data[!(names(train_data)) %in% c("
    ViolentCrimesPerPop")],
3                       test = NULL,
4                       train_data$ViolentCrimesPerPop, k = 1,
5                       algorithm = c("kd_tree"))
6 plot(nn_regr$pred, nn_regr$residuals)
7
8 # Evaluate the regression on the above training data with mean-squared
    error.
9 mse_residuals_knn <- sum((nn_regr$residuals - mean(nn_regr$residuals))
    ^ 2) /
10   length(nn_regr$residuals)
11 printf(" - Mean-squared error on the training data (80%%): %.2e", mse_
    residuals_knn)

```

```

12
13 # Compute the regression on the above test data.
14 klabels <- knn(train_data[!(names(train_data)) %in% c("
      ViolentCrimesPerPop")],
15               test_data [!(names(test_data)) %in% c("
      ViolentCrimesPerPop")],
16               train_data$ViolentCrimesPerPop, k = 1,
17               prob = FALSE, algorithm = c("kd_tree"))
18 k_test_residuals <- test_data$ViolentCrimesPerPop - as.numeric(levels(
      klabels))[klabels]
19 plot(as.numeric(levels(klabels))[klabels], k_test_residuals)
20
21 # Evaluate the previous regression with mean-squared error and print
    the result.
22 mse_residuals_knn_test <- sum((k_test_residuals - mean(k_test_residuals
      )) ^ 2) /
23   length(k_test_residuals)
24 printf(" - Mean-squared error on the test data (20%%): %.2e", mse_
      residuals_knn_test)

```

3.2 Results

3.3 Conclusions

4 Dealing with missing values

4.1 Implementation

```

1 printf(" - Mean-squared error on the test data (20%%): %.2e", mse_
      residuals_knn_test)
2
3 # Impute unknown values by covering all the columns with unknowns.
4 interpolated_data <- clean_data
5 for(i in attr_with_unkw) {
6   # Divide data by rows in two sets: one with the samples that
      contain a question
7   # mark in that column and one with the rest.
8   unk_indices <- which(clean_data[,i] == '??')
9   i_unk_data <- no_unkw_data[unk_indices, ]
10  i_no_unk_data <- no_unkw_data[-unk_indices, ]
11
12  # Compute the nearest neighbours in the set of elements with
      question mark of
13  # the elements with a question mark in that column.
14  i_klabels <- knn(i_no_unk_data,
15                  i_unk_data,
16                  i_no_unk_data$ViolentCrimesPerPop, k = 1,
17                  prob = FALSE, algorithm=c("kd_tree"))
18  indices <- attr(i_klabels, "nn.index")
19  nn_subs <- clean_data[-unk_indices, ][indices,]
20
21  # Substitute the question mark values by the value of the nearest
      neighbour.

```

```

22     interpolated_data[unk_indices, i] <- nn_subs[,i]
23 }
24
25 # Split off new test data.
26 folds <- cvFolds(nrow(interpolated_data), K=5)
27 train_data <- interpolated_data[folds$subsets[folds$which != 1], ]
28 test_data <- interpolated_data[folds$subsets[folds$which == 1], ]
29
30 # Compute a linear regression again and evaluate it with mean-squares.
31 printf("Linear Regression (imputed missing values):")
32 linear_regr_test <- lm(ViolentCrimesPerPop ~ ., data=train_data)
33
34 residuals_test <- test_data$ViolentCrimesPerPop - predict(linear_regr_
    test, test_data)
35 plot(test_data$ViolentCrimesPerPop, residuals_test)
36
37 mse_residuals_test <- sum((residuals_test - mean(residuals_test)) ^ 2)
    / length(residuals_test)
38 printf(" - Mean-squared error on the test data (20%%): %.2e", mse_
    residuals_test)
39
40 # Also, compute a nearest neighbour regression and evaluate it.
41 printf("Nearest Neighbours (imputed missing values):")
42 klabels <- knn(train_data[!(names(train_data)) %in% c("
    ViolentCrimesPerPop")],
43               test_data [!(names(test_data)) %in% c("
    ViolentCrimesPerPop")],
44               train_data$ViolentCrimesPerPop, k = 1,
45               prob = FALSE, algorithm = c("kd_tree"))
46 k_test_residuals <- test_data$ViolentCrimesPerPop - as.numeric(levels(
    klabels))[klabels]
47 plot(as.numeric(levels(klabels))[klabels], k_test_residuals)
48
49 # Compute the mean-squared error and print the result.
50 mse_residuals_knn_test <- sum((k_test_residuals - mean(k_test_residuals
    )) ^ 2) /
51     length(k_test_residuals)
52 printf(" - Mean-squared error on the test data (20%%): %.2e", mse_
    residuals_knn_test)

```

4.2 Results

4.3 Conclusions

5 Modified Nearest Neighbours

5.1 Implementation

5.2 Results

5.3 Conclusions